

Responsible AI on Monday

Session length: 30 minutes

Core practice: Use AI to support a task, then complete a deliberate human review before using the output.

1. Who this is for

Academic and professional services staff who:

- have not read the report and are unlikely to read it;
- use, or may soon use, generative AI in teaching, assessment, research or administration;
- need one responsible practice they can apply in their own work on Monday;
- do not need technical AI knowledge.

The session focuses on the report's central principle: AI may support work, but responsibility for accuracy, decisions, ethics and professional conduct remains with people. [\[Deep Resea...dByCopilot | PDF\]](#)

Time: 2 minutes

2. Objectives

By the end of the session, when given a realistic AI-assisted task, participants will be able to:

1. **Identify** one risk that requires human attention before an AI output is used.
2. **Apply** a four-question human review to an AI-generated output.
3. **Decide** whether to use, revise or reject that output.
4. **Write** one responsible AI action they will use in their own work on Monday.

These objectives use observable actions rather than “understand” or “be aware of”, because those internal states cannot be seen or checked.

Time: 2 minutes

3. What the document actually says

The sector has moved beyond “should we use AI?”

The report says generative AI had become part of normal academic practice by 2026. It presents the main question as how universities engage with it safely, fairly and effectively, rather than whether they should engage at all. It reports that 95% of UK undergraduates use

AI in some way and 94% use generative AI to support assessed work. [\[Deep Resea...dByCopilot | PDF\]](#)

Four features of stronger implementation

The report identifies four characteristics:

1. secure institutional AI infrastructure;
2. assessment redesign rather than reliance on detection;
3. capability building rather than policy alone;
4. preserved human accountability. [\[Deep Resea...dByCopilot | PDF\]](#)

Risks apply to different kinds of work

Risk from the report	What the report says
Data privacy and intellectual property	Sensitive information may be exposed through inappropriate use of public AI systems.
Digital equity	Unequal access may widen existing educational disadvantage.
Cognitive over-reliance	People may outsource reading, writing or analysis rather than developing capability.
AI detection failure	Detection approaches remain unreliable for high-stakes decisions.
Environmental impact	Large-scale AI use consumes substantial computing resources.

[\[Deep Resea...dByCopilot | PDF\]](#)

The common thread: accountable human judgement

The report repeatedly assigns the final role to a person:

- academics retain final judgement over AI-supported feedback;
- domain experts verify extracted research information;
- professional judgement remains central to assessment;
- staff remain accountable when AI supports administrative work. [\[Deep Resea...dByCopilot | PDF\]](#)

Facilitator line:

“You do not delegate accountability when you delegate part of a task to AI. Your work represents you.”

Time: 6 minutes

4. Walkthrough

Worked example: using AI to summarise staff feedback

Situation

A professional services colleague asks an approved AI tool to summarise comments collected after a staff event. The AI produces four themes and a short recommendation for the next event.

Before copying the summary into a briefing note, the colleague applies a human review.

The four-question human review

1. Should this information be entered?

Ask:

- Does the material contain sensitive, confidential or personal information?
- Am I using an appropriate institutional environment?

The report warns that inappropriate use of public AI systems may expose sensitive information. It recommends enterprise platforms, staff training and suitable data-retention arrangements as mitigation. [\[Deep Resea...dByCopilot | PDF\]](#)

Decision in this example: Remove names and identifiers before processing. Use the institutionally approved system.

2. Is the output supported?

Ask:

- Can I trace each theme back to the original comments?
- Has the AI introduced a claim that respondents did not make?

The report says institutional platforms cannot verify every fact or detect every hallucination. It also says staff may accept inaccurate outputs without sufficient checking. [\[Deep Resea...dByCopilot | PDF\]](#)

Decision in this example: Check every proposed theme against the source comments. Delete an unsupported claim.

3. Is the output fair and appropriate?

Ask:

- Has the summary hidden minority or critical views?
- Would unequal access, context or representation alter the conclusion?
- Is the language suitable for the audience?

The report identifies digital equity as a strategic risk and states that people provide contextual understanding, empathy and discretion. [\[Deep Resea...dByCopilot | PDF\]](#)

Decision in this example: Add an important concern raised by a small number of respondents rather than presenting only the majority view.

4. *Will I take responsibility for it?*

Ask:

- Would I be comfortable explaining how this was produced?
- Can I defend the final wording and decision?
- Should I use it, revise it or reject it?

The report states that responsibility for decisions, accuracy, ethics and professional conduct must remain with people. [\[Deep Resea...dByCopilot | PDF\]](#)

Final decision: Revise the summary, record that AI supported the analysis, and have the responsible colleague approve the briefing.

The resulting practice

Before using an AI output: check the input, verify the output, consider the people affected, and own the final decision.

Time: 7 minutes

5. Practice

Activity: Monday Human Review

Individual work: 3 minutes

Choose one task you might use AI for on Monday:

- drafting an email or document;
- summarising information;
- analysing feedback;
- preparing teaching material;
- creating draft feedback;
- retrieving or organising information.

Write brief answers to these questions:

1. **Input:** What information must I avoid entering, remove or protect?
2. **Evidence:** What source will I use to check the output?
3. **Impact:** Who could be affected if the output is incomplete, inaccurate or unfair?
4. **Accountability:** What must I personally decide before the output is used?
5. **Action:** Will I use, revise or reject it?

Pair discussion: 2 minutes

Explain your planned review to another participant. The other person identifies one check you may have missed.

What a strong response contains

A strong response:

- names a real task;
- identifies a realistic risk;
- states what evidence will be checked;
- identifies who may be affected;
- leaves the final decision with a named role or person;
- does not treat polished AI wording as proof of accuracy.

These quality criteria are the training activity's **house convention**. The report supports human verification and accountability, but it does not prescribe this exact five-question activity. [\[Deep Resea...dByCopilot | PDF\]](#)

Time: 5 minutes

6. Check your understanding

Scenario 1: AI-generated student feedback

An academic uses AI to generate draft feedback aligned with a marking rubric. The wording is clear and consistent.

What should happen next?

- A.** Send it because the rubric was included in the prompt.
- B.** Check it against the student's work and rubric, revise it where needed, and retain final academic judgement.
- C.** Reject it because AI must not contribute to academic feedback.

Correct answer: B

- **Why B is right:** The report describes AI-generated draft feedback aligned with rubrics, but final judgement remains with the academic marker.
- **Why A is wrong:** A prompt and polished response do not remove the risk of error or over-reliance.
- **Why C is wrong:** The report includes supervised AI-supported feedback as a practical action. It does not call for a blanket prohibition. [\[Deep Resea...dByCopilot | PDF\]](#)

Scenario 2: AI meeting summary

A colleague receives an AI-generated summary of a meeting. One action sounds plausible, but they cannot remember anyone agreeing to it.

What should they do?

- A. Include it because the AI probably captured something they missed.
- B. Remove the entire summary because AI outputs cannot be used.
- C. Check the action against the available meeting evidence and confirm it with the responsible person before treating it as agreed.

Correct answer: C

- **Why C is right:** The report supports AI for meeting analysis but says that human accountability must be retained and inaccurate outputs require checking.
- **Why A is wrong:** Plausibility is not evidence.
- **Why B is wrong:** The report supports supervised workflow augmentation rather than rejecting AI-supported administration. [\[Deep Resea...dByCopilot | PDF\]](#)

Time: 5 minutes

7. Where to go next

Monday commitment

Ask each participant to complete this sentence:

On Monday, when I use AI to _____, I will check _____ before I use the output.

Examples:

- “When I use AI to structure teaching material, I will check every factual claim against the original sources.”
- “When I use AI to summarise feedback, I will check that minority views have not disappeared.”
- “When I use AI to draft an email, I will check the content, tone and recipients before sending.”
- “When I use AI to generate draft feedback, I will compare it with the work and make the final judgement myself.”

Close with:

AI can support the work. It cannot carry your accountability. Your work represents you.

Time: 3 minutes

Duration check

Section	Time
Who this is for	2 minutes
Objectives	2 minutes
What the document actually says	6 minutes
Walkthrough	7 minutes
Practice	5 minutes
Check your understanding	5 minutes
Where to go next	3 minutes
Total	30 minutes

Gap note

The source is silent or incomplete in the following areas:

- **Approved local AI systems:** NONE STATED.
- **Cardiff University rules for entering personal, confidential, research or student information:** NONE STATED.
- **Local route for reporting an AI-related data or information incident:** NONE STATED.
- **Required form of disclosure when staff use AI in teaching, research or professional services:** NONE STATED.
- **Named institutional owner for responsible AI implementation:** NONE STATED.
- **Required verification standard for different levels of risk:** NONE STATED.
- **Activities that must not use AI:** NONE STATED.
- **Whether a second person must review high-impact outputs:** NONE STATED.
- **Local accessibility requirements for AI-generated material:** NONE STATED.
- **The four-question review used in this session:** House convention derived from the report's risks and human-accountability principle, not a framework specified by the report.
- **The use, revise or reject decision:** House convention, not terminology specified by the report.
- **The Monday commitment sentence:** Training device, not a requirement from the report.
- **The report's evidence base has limits:** its own examiner's note scores sourcing at 2/5 and evidence versus opinion, recency, learning and examples at 3/5. Some case studies cite an earlier deep-research PDF rather than a named primary source. [\[Deep Resea...dByCopilot | PDF\]](#)

a person could replace the generic walkthrough with one approved, real Cardiff University case and add the relevant local policy, system and escalation route.